

# R DATA WRANGLING

## LECTURE 12

---

Dirk Eddelbuettel

STAT 430: Data Science Programming Methods (Fall 2019)

Department of Statistics, University of Illinois

# DATA WRANGLING

---

## Overview

- In the last three lecture we covered
  - R data types
  - R control flow and functions
  - R data input/output
- So now let us learn how to slice + dice + aggregate some data
  - This lecture focuses on Base R and functions come with R
  - You can assume that these will always be at your disposal
- A solid grasp of basics will aid in using alternate approaches
  - `data.table` (next lecture)
  - `dplyr` (lecture after next)

### Three key aspect we cover

- Alter / expand the data set by computing new columns
- Summarize / describe the data, also by taking simple subsets
- Conditional summaries based on some columns

## Flights Data Set

```
> ## A large flights data set we can read in directly
> ## Also present in data-examples repo
> url <- paste0("https://github.com/Rdatatable/",
+              "data.table/blob/master/vignettes/",
+              "flights14.csv?raw=true")
> data <- read.csv(url)
> dim(data)
# [1] 253316      11
```

Define and assign

```
> ## using $ for columns  
> data$tot_delay <- data$dep_delay + data$arr_delay  
> ##  
> ## equivalent  
> data[, "tot_delay"] <- data[, "dep_delay"] +  
+     data[, "arr_delay"]
```

Generally any *expression* can be assigned to a new column (which will be added), or equally to an existing column that will be overwritten.

## Use `within`

Because the explicit enumeration as in `data[, "dat_delay"]` etc is cumbersome, a more functional alternative exists:

```
> data <- within(data,  
+               tot_delay <- dep_delay + arr_delay)
```

## Use with

Similar, `with` can give more “direct” access to the column identifiers:

```
> with(data, mean(distance))  
# [1] 1099.445  
> with(data, mean(distance)) == mean(data[, "distance"])  
# [1] TRUE
```



## Another with example

```
> ## nicer expression using with ...  
> with(mtcars, mpg[cyl == 8 & disp > 350])  
# [1] 18.7 14.3 10.4 10.4 14.7 19.2 15.8  
> ##  
> ## ... than explicitly referencing  
> mtcars$mpg[mtcars$cyl == 8 & mtcars$disp > 350]  
# [1] 18.7 14.3 10.4 10.4 14.7 19.2 15.8
```

The `transform` function can be used as well:

```
> data <- transform(data,  
+                   total_delay = arr_delay+dep_delay)
```

Note that we use `=` for assignment with in the scope of the `data` object here.

## Useful to “sweep” a function across

- The **apply** function can operate on matrices and data.frames
- It takes three arguments:
  - the object
  - the direction: 1 for row-wise, 2 for column-wise
  - a function

## Example

```
> M <- matrix(1:9,3,3)
> M
#      [,1] [,2] [,3]
# [1,]    1    4    7
# [2,]    2    5    8
# [3,]    3    6    9
> apply(M, 1, sum) # rows
# [1] 12 15 18
> apply(M, 2, sum) # cols
# [1]  6 15 24
```

```
> M <- matrix(1:9,3,3)
> D <- data.frame(M); D
#   X1 X2 X3
# 1  1  4  7
# 2  2  5  8
# 3  3  6  9
> apply(D, 1, sum) # rows
# [1] 12 15 18
> apply(D, 2, sum) # cols
# X1 X2 X3
#  6 15 24
```

# SUBSETS

---

## Index Expression

We can use any valid expression for the row indices

```
> mean(data[data$origin=="JFK" &  
+         data$dest=="LAX", "tot_delay"])  
# [1] 11.66105
```

This works because the expression `data$origin=="JFK" & data$dest=="LAX"` returns a logical vector of the same size as `data` so it can be used for indexing.

(But let's recall that purely numerical indexing also work.)

We can save one layer of `$` subsetting via `with`:

```
> with(data, mean(data[origin=="JFK" &  
+                  dest=="LAX", "tot_delay"]))  
# [1] 11.66105
```

## SUMMARIES BY GROUP

---



Using factor variables provides logical grouping, especially when the cardinality of the factor variable is low. *I.e.* here we would rather condition on **origin** with its three values than **dest** which has 109.

### **aggregate**

In its simplest form, **aggregate** takes a vector argument, followed by a list of conditioning variable, and a function:

```
> aggregate(data$arr_delay,  
+           list(origin=data$origin), mean)  
#   origin      x  
# 1    EWR 10.026121  
# 2    JFK  7.731465  
# 3    LGA  6.601968
```

## aggregate

The `aggregate` function also has a more powerful formula interface

```
> ## average arrival delay by airport
> aggregate(arr_delay ~ origin, data, mean)
#   origin arr_delay
# 1    EWR 10.026121
# 2    JFK  7.731465
# 3    LGA  6.601968
```

Use `cbind` to select several columns for multiple summaries

```
> aggregate(cbind(arr_delay, dep_delay) ~ origin,  
+           data, mean)  
#   origin arr_delay dep_delay  
# 1    EWR 10.026121  15.21248  
# 2    JFK  7.731465  11.44617  
# 3    LGA  6.601968  10.60500
```

The right-hand side of the formula can expand to provide multiple groups:

```
> head(aggregate(tot_delay ~ origin + carrier,  
+               data, mean), 7)  
#   origin carrier tot_delay  
# 1    EWR      AA 24.597206  
# 2    JFK      AA 15.938271  
# 3    LGA      AA  9.578602  
# 4    EWR      AS  4.942509  
# 5    EWR      B6 18.480906  
# 6    JFK      B6 22.785535  
# 7    LGA      B6 22.019850
```

## Summary of Formula notation

term	example	description
~	$y \sim x$	model $y$ as a function of $x$
+	$y \sim a + b$	include columns $a$ and $b$
-	$y \sim a - b$	include $a$ but not $b$
:	$y \sim a : B$	estimate interaction of $a$ and $b$
*	$y \sim a * b$	include $a$ and $b$ and their interaction
	$y \sim a   b$	estimate $y$ as function of $a$ conditional on $b$

This will become more relevant when doing modeling.

Source: Table 13-3 in de Vries and Meys (2012)

## tapply

Similar to aggregate

```
> tapply(data$tot_delay, data$origin, mean)
#      EWR      JFK      LGA
# 25.23860 19.17763 17.20697
```

tapply: Multiple indices

```
> head(tapply(data$tot_delay,  
+           list(data$carrier, data$origin), mean), 8)  
#           EWR           JFK           LGA  
# AA 24.597206 15.93827  9.578602  
# AS  4.942509           NA           NA  
# B6 18.480906 22.78553 22.019850  
# DL 25.699254 16.51813 16.366738  
# EV 31.597748 36.90833 27.939167  
# F9           NA           NA 51.321353  
# FL           NA           NA 34.269384  
# HA           NA 20.93846           NA
```

Note how **tapply** preserves all cells in the table.

by: aggregate by factor

```
> by(data$arr_delay, data$origin, summary)
```

```
# data$origin: EWR
```

#	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
#	-71.00	-14.00	-3.00	10.03	17.00	1494.00

```
# -----
```

```
# data$origin: JFK
```

#	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
#	-79.000	-15.000	-3.000	7.731	14.000	1223.000

```
# -----
```

```
# data$origin: LGA
```

#	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
#	-112.000	-16.000	-5.000	6.602	14.000	996.000



# THE LAPPLY/SAPPLY/MAPPLY/... FUNCTIONS

---

## Loop over list

- `lapply` is very powerful and useful
  - it takes a list,
  - then applies then given function to each element
  - and returns a list of results
- *internally* a `data.frame` *is* a list so we can use it
- `lapply` is very powerful and useful
- `do.call()` can then be used to compact the results

## Simple example

```
> set.seed(123) # make it reproducible
> alist <- list(A=matrix(1:9,3),B=seq(20,30),C=rnorm(5))
> lapply(alist, sum)
# $A
# [1] 45
#
# $B
# [1] 275
#
# $C
# [1] 0.9678513
```

## A simplified version

- `sapply` is like `lapply`
- but returns a simpler object, where possible
- if the operation reduces to a number we get a vector back

## Simple example

```
> set.seed(123) # make it reproducible
> alist <- list(A=matrix(1:9,3),B=seq(20,30),C=rnorm(5))
> sapply(alist, sum)
#           A           B           C
# 45.0000000 275.0000000  0.9678513
```

## A multivariate **sapply**

- it sweeps the function over the  $i$ -th elements of *multiple* lists
- can often be used where a loop might be needed

```
> mapply(paste0, LETTERS[1:5], letters[1:5])  
#      A      B      C      D      E  
# "Aa" "Bb" "Cc" "Dd" "Ee"
```

## Another data.frame operation

- The `merge()` function combines data.frame objects
- There is a fairly close connection to SQL's **JOIN** command
- Try the examples on the right
- See `help(merge)` for more

```
A <- data.frame(x=c("A","C","E"),
                y=c(1,3,5))
B <- data.frame(x=c("B", "C", "D"),
                y=c(20,30,40))

merge(A, B, by="x")           # inner
merge(A, B, by="x", all=TRUE) # outer
merge(A, B, by="x", all.x=TRUE) # left
merge(A, B, by="x", all.y=TRUE) # right
```

## Map, Reduce, Filter

- R also has a set of function supporting a functional approach
- **Map(f, x)** applies binary **f()** to each element of **x**
  - this returns a list
  - try *e.g.* **Map(sqrt, 1:5)**
- **Reduce(f, x)** applies **f(x,y)** over successive **x** elements
  - this (typically) returns a scalar
  - try *e.g.* **Reduce(sum, 1:5)**
- **Filter(f, x)** select elements where the predicate is true
- and more – we won't be using (or testing) these
- but it is useful to know about them



## Data Wrangling

We have seen how

- to alter a **data.frame** by adding to it
- to operate on explicit subsets via indexing
- to operate on implicit groupings from factor
- to use some of the **\*apply** functions
- to combine **data.frame** object via **merge**
- to use functional programming over lists and **data.frame** objects

These commands are versatile but at times a little cumbersome and idiosyncratic. They are however universal as part of R.

We will see alternatives for basic wrangling in the next two lectures.