

INTRO TO PARALLEL COMPUTING AND BIG(GER) DATA

LECTURE 15

Dirk Eddebuettel

STAT 430: Data Science Programming Methods (Fall 2019)

Department of Statistics, University of Illinois

External reads

- [Shorter tutorial](#) on *Parallel Computing in R* by Matt Jones
- [Comprehensive tutorial](#) by Jonathan Dursi (with [sources](#))
- [Parallel Computing for Data Science](#) book by Norm Matloff
- CRAN Task View on [High-Performance Computing](#)

R: Our engine

- R itself is “single-threaded”
- That means a single *thread* of execution in our program
- (There are some exceptions but this broadly holds)
- The running program is generally referred to as a process
- The internal design of R makes this hard / impossible to change
- So this puts a limit on how much work we can do “at once”

Yet ...

- Computers have become multi-core
- Even laptops and cell phone have four or more cores
- Servers can have dozens of cores (and more than one cpu)
- And then we can of course have multiple servers
- Two broad answers:
 - *implicit* parallel work on the same computer (our focus here)
 - *explicit* parallel work across different computers (not here)

Parallel work on the same computer

- Our running process may *fork* other processes
- Our (single-threaded) process may contain multi-threaded parts
 - generally well-shielded subroutines with well-defined entry
 - R itself has a few multithreaded helper functions
 - **data.table** is an example: many parallelised routines
 - this often involved multithreaded C/C++ code
 - using either the **OpenMP** or **phreads** libraries
 - and that is not our focus here
- So focus will be on R tools mostly for multi-process work
- Some platform-dependency, see **parallel** package vignette

One simple exception

- R can be set up to use parallel Matrix math libraries
- This depends on how R is built / compiled
- You need to check the details on your installation
- The `sessionInfo()` function displays it
- Look for BLAS (== Basic Linear Algebra Subroutine) and LAPACK (== Linear Algebra PACKage)

On my computer here:

```
R> sessionInfo()
```

```
R version 3.6.1 (2019-07-05)
```

```
Platform: x86_64-pc-linux-gnu (64-bit)
```

```
Running under: Ubuntu 19.04
```

```
Matrix products: default
```

```
BLAS: /usr/lib/x86_64-linux-gnu/openblas/libblas.so.3
```

```
LAPACK: /usr/lib/x86_64-linux-gnu/libopenblas-p-r0.3.5.so
```

Here OpenBLAS signals that such a parallel BLAS/LAPACK implementation which in fact multithreaded so e.g. a matrix multiplication can multiply and some several row/column combinations at one.

On our RStudio Cloud instance:

```
> sessionInfo()  
R version 3.5.3 (2019-03-11)  
Platform: x86_64-pc-linux-gnu (64-bit)  
Running under: Ubuntu 16.04.6 LTS
```

```
Matrix products: default  
BLAS: /usr/lib/atlas-base/atlas/libblas.so.3.0  
LAPACK: /usr/lib/atlas-base/atlas/liblapack.so.3.0
```

This is also a decent choice: ATLAS stands for Automatically Tuned Linear Algebra Subroutines – optimized, but *not* multithreaded (as the cloud server will likely already be part of a shared larger computer).

Key Points

- **parallel** is included with R, part of every installation
- Consolidates parts of earlier packages
 - **multicore** which offered easy use of multiple cores
 - **snow** which combines both *implicit* and *explicit* parellism
 - **snow** is still available as an external package
- pdf vignette for package **parallel** is a very good introduction

Two key functions

- `parLapply(cl, x, FUN, ...)` as a parallel `lapply`
 - uses a `cl <- makeCluster(...args...)` as first argument
 - and has to call `stopCluster(cl)`
- `mclapply(X, FUN, ..., mc.cores)` ditto
 - just dispatches to the `mc.cores`
 - no need for `cl` object (easier, less flexible)
 - but `mclapply` is not available on Windows
- uses process parallelism either way

Easy and reliable work-horse

- generalizes `lapply`
- sweeps a function over all elements of a vector or list
- easy to use on single Linux or macOS machine

- Background: every process on Unix has a process id, or 'pid'
- Tools like **ps** or **top** show it
- R wraps many system functions, here we can use `Sys.getpid()`

```
Sys.getpid()
```

```
## [1] 2601
```

The displayed number will vary, but if you **ps -vax** and **grep** for the same id (on the shell) you should see your R process.

- Now let's run this six times via `lapply`
- We immediately collapse the result into a vector for compactness
- The function used by `lapply` must take an argument so we wrap `pid()`

```
do.call(c, lapply(1:6, function(x) Sys.getpid()))
```

```
## [1] 2601 2601 2601 2601 2601 2601
```

Six times the same value. Do you know why?

- Now let's run this six times via `mclapply`
- We immediately collapse the result into a vector for compactness

```
suppressMessages(library(parallel))  
do.call(c, mclapply(1:6, function(x) Sys.getpid()))
```

```
## [1] 2649 2650 2651 2652 2649 2650
```

Four different values, two repeats. Can you think of a reason why?

- We just illustrated process-parallel operations
- `lapply` ran sequentially in the same R process:
 - same process id each time
 - (though random across different R sessions)
- `mclapply` spawned different processes
 - different process id values

```
suppressMessages(library(boot))
cd4.rg <- function(data, mle) MASS::mvrnorm(nrow(data), mle$m, mle$v)
cd4.mle <- list(m = colMeans(cd4), v = var(cd4))
cd4.boot <- boot(cd4, corr, R = 999, sim = "parametric",
                ran.gen = cd4.rg, mle = cd4.mle)
boot.ci(cd4.boot, type = c("norm", "basic", "perc"),
        conf = 0.9, h = atanh, hinv = tanh)
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
```

```
## Based on 999 bootstrap replicates
```

```
##
```

```
## CALL :
```

```
## boot.ci(boot.out = cd4.boot, conf = 0.9, type = c("norm", "basic",
```

```
##      "perc"), h = atanh, hinv = tanh)
```

```
##
```

```
## Intervals :
```

```
## Level      Normal      Basic      Percentile
```

```
## 90% ( 0.4528, 0.8593 ) ( 0.4565, 0.8644 ) ( 0.4761, 0.8706 )
```

```
## Calculations on Transformed Scale; Intervals on Original Scale
```



```
suppressMessages(library(parallel))
mc <- 4 # adjusted from 2, also adjusted R on next line
run1 <- function(...) boot(cd4, corr, R = 250, sim = "parametric",
                           ran.gen = cd4.rg, mle = cd4.mle)
set.seed(123, "L'Ecuyer") # To make this reproducible
cd4.boot <- do.call(c, mclapply(seq_len(mc), run1) )
boot.ci(cd4.boot, type = c("norm", "basic", "perc"),
        conf = 0.9, h = atanh, hinv = tanh)

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = cd4.boot, conf = 0.9, type = c("norm", "basic",
##          "perc"), h = atanh, hinv = tanh)
##
## Intervals :
## Level      Normal              Basic              Percentile
## 90%   ( 0.4569, 0.8553 )   ( 0.4549, 0.8565 )   ( 0.4995, 0.8711 )
## Calculations on Transformed Scale; Intervals on Original Scale
```

Creating two functions:

```
bstrpSer <- function() {  
  cd4.boot <- boot(cd4, corr, R = 999, sim = "parametric",  
                  ran.gen = cd4.rg, mle = cd4.mle)  
  boot.ci(cd4.boot, type = c("norm", "basic", "perc"),  
          conf = 0.9, h = atanh, hinv = tanh)  
}  
  
bstrpPar <- function() {  
  run1 <- function(...) boot(cd4, corr, R = 250, sim = "parametric",  
                             ran.gen = cd4.rg, mle = cd4.mle)  
  cd4.boot <- do.call(c, mclapply(seq_len(mc), run1) )  
  boot.ci(cd4.boot, type = c("norm", "basic", "perc"),  
          conf = 0.9, h = atanh, hinv = tanh)  
}
```

And timing them:

```
library(rbenchmark)  
benchmark(bstrpSer(), bstrpPar(), replications=100)[,1:4]
```

##	test	replications	elapsed	relative
## 2	bstrpPar()	100	3.620	1.000
## 1	bstrpSer()	100	6.084	1.681

so here we see a 61% timing difference on the simple example.

Important lesson: speedup from parallel task almost always sublinear as overhead involved.

foreach

- The **foreach** package is popular and widely used
- It uses an ‘iterator’ approach to ‘loop’ over a set of loop indices
- Different backends can be registered
- This makes it easy to switch between serial and parallel modes

Standard **for** loop

```
for (i in 1:3) print(sqrt(i))
```

```
## [1] 1  
## [1] 1.41421  
## [1] 1.73205
```

Using **foreach** in serial mode

```
library(foreach)  
rl <- foreach (i=1:3) %do% sqrt(i)  
do.call(c, rl) # list into vector
```

```
## [1] 1.00000 1.41421 1.73205
```

- **%do%** operates in the object created by **foreach**
- a `{ ... }` block can follow as well as single call seen here

foreach and %doParallel%

- Several different backends available
- One of the simplest is **doParallel**
- We register a backend, and use **%dopar** instead of **%do**

foreach and %doParallel%

```
suppressMessages(library(doParallel))  
cl <- makeCluster(4)  
registerDoParallel(cl) # use multicore-style forking  
rl <- foreach (i=1:3) %dopar% sqrt(i)  
do.call(c, rl)
```

```
## [1] 1.00000 1.41421 1.73205
```

```
stopCluster(cl) # cleanup is good practice
```

- New(er) package (family), may replace **foreach** use over time
- Many interesting features but also easy *parallel* use:

```
# two functions from package 'future'
plan(multiprocess)
# function from package 'future.apply'
rl <- future_lapply(countries, FUN=doWork)
res <- do.call(rbind, rl) # combine into single list
```

Registers for multiprocess use (on same machine), applies **doWork** across counties and combines results

NOTES

General comments

- We want jobs 'large enough' to be worthwhile running in parallel
- Overhead in coordinating runs so using 2,3,4,... threads or cores or machines generally does *not* scale 2,3,4,... times
- Many tasks are 'embarrassingly parallel':
 - Bootstrap and Monte Carlos simulation
 - Markov Chain Monte Carlo and Gibbs sampling
 - Machine Learning calibration of hyper-parameters
 - Machine Learning crossvalidation across distinct data sets
 - Data parallelism: different days or regions or sources ...

General comments

- Key is not to have dependence between data sets
- `mclapply()` easy to use and great workhorse
- `future` package (and `future.apply` etc) another (newer) excellent candidate

R prefers data in its memory

- That is a constraint and “feature” for exploratory analysis
- But because of “big data” alternatives exist (see next slide)
- A lot of this is too specific and requires special setup
- Just know it is there so you can catch up when the need arises

R Packages that help

- **biglm** can run regressions etc out of memory
- **bigmemory** keeps data in RAM but outside of R (this avoids extra copies)
- additional packages build on top of **bigmemory** for analysis
- **disk.frame** uses **fst** to treat data as it were local
- computation inside the database (*i.e.* **dplyr** and **dbplyr**)
- and of course R as part of large batch systems (Spark, Hadoop)
- a bit more on Big Data is at this [CRAN Task View](#)

Simple example using `iris`

```
library(fst)
# generate a sample fst file
path <- paste0(tempfile(), ".fst")
# write iris data to file
write_fst(iris, path)
```

```
ft <- fst(path)      # create a fst_table object that can be used as a data frame
print(ft)            # print snapshot view
```

```
## <fst file>
## 150 rows, 5 columns (file16b86eb661d8.fst)
##
##      Sepal.Length Sepal.Width Petal.Length Petal.Width  Species
##      <double>      <double>      <double>      <double> <factor>
## 1           5.1           3.5           1.4           0.2    setosa
## 2           4.9           3.0           1.4           0.2    setosa
## 3           4.7           3.2           1.3           0.2    setosa
## 4           4.6           3.1           1.5           0.2    setosa
## 5           5.0           3.6           1.4           0.2    setosa
## --           --           --           --           --      --
## 146          6.7           3.0           5.2           2.3  virginica
## 147          6.3           2.5           5.0           1.9  virginica
## 148          6.5           3.0           5.2           2.0  virginica
## 149          6.2           3.4           5.4           2.3  virginica
## 150          5.9           3.0           5.1           1.8  virginica
```

Some more code examples – not executed, see `example(fst)`

```
# select columns and rows
x <- ft[10:14, c("Petal.Width", "Species")]

# use the common list interface as with data.frame
ft[TRUE]
ft[c(TRUE, FALSE)]
ft[["Sepal.Length"]]
ft$Petal.Length

# use data frame generics
nrow(ft)
dim(ft)
colnames(ft)
rownames(ft)
```


SUMMARY

Parallel Programming in R

- Multithreading at C/C++ source level (for experts)
- Multiprocess via packages like `parallel` and `future`
- Switching from `lapply` to `mclapply` is very easy
- `future` package best path forward

Bigger Data

- Several solutions via packages
- New package `fst` looks very promising